

Data Capstone-Healthcare

July 10, 2023

1 Capstone Project: Healthcare

Problem Statement: * NIDDK (National Institute of Diabetes and Digestive and Kidney Diseases) research creates knowledge about and treatments for the most chronic, costly, and consequential diseases. * The dataset used in this project is originally from NIDDK. The objective is to predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. * Build a model to accurately predict whether the patients in the dataset have diabetes or not.

Dataset Description: The datasets consists of several medical predictor variables and one target variable (Outcome). Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and more.

2 Solution:

2.1 Week 1:

2.1.1 Data Exploration:

(1) Read Data and Perform descriptive analysis:

```
[1]: # Importing required libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
from warnings import filterwarnings
filterwarnings('ignore')
%matplotlib inline
```

```
[2]: #reading the dataset
df=pd.read_csv('diabetes.csv')
```

```
df.head()
```

```
[2]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0         6      148         72           35          0  33.6
1         1       85         66           29          0  26.6
2         8     183         64           0           0  23.3
3         1       89         66           23         94  28.1
4         0     137         40           35        168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
4                2.288    33         1
```

```
[ ]:
```

```
[3]: #shape of the dataset
df.shape
```

```
[3]: (768, 9)
```

```
[4]: # checking for missing values
df.isnull().sum()
```

```
[4]: Pregnancies      0
Glucose              0
BloodPressure        0
SkinThickness        0
Insulin              0
BMI                  0
DiabetesPedigreeFunction  0
Age                  0
Outcome              0
dtype: int64
```

```
[5]: # checking duplicate rows in dataset
df.duplicated().sum()
```

```
[5]: 0
```

```
[6]: # getting information about dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
```

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

```
[7]: # getting statical information about our dataset
df.describe().transpose()
```

```
[7]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

(2) Visually explore these variables using histograms. Treat the missing values accordingly.

```
[8]: df[df['Glucose']==0]
```

```
[8]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
75	1	0	48	20	0	24.7	
182	1	0	74	20	23	27.7	


342	1	0	68	35	0	32.0
349	5	0	80	32	0	41.0
502	6	0	68	41	0	39.0

	DiabetesPedigreeFunction	Age	Outcome
75	0.140	22	0
182	0.299	21	0
342	0.389	22	0
349	0.346	37	1
502	0.727	41	1

Visually explore these variables using histograms. Treat the missing values accordingly.

```
[9]: plt.hist(df['Glucose'])
```

```
[9]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
      array([ 0. , 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
            179.1, 199. ]),
      <BarContainer object of 10 artists>)
```



```
[10]: df[df['BloodPressure']==0].head()
```


```
[10]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
7          10     115           0           0           0   35.3
15          7     100           0           0           0   30.0
49          7     105           0           0           0    0.0
```

60	2	84	0	0	0	0.0
78	0	131	0	0	0	43.2

	DiabetesPedigreeFunction	Age	Outcome
7	0.134	29	0
15	0.484	32	1
49	0.305	24	0
60	0.304	21	0
78	0.270	26	1

```
[11]: plt.hist(df['BloodPressure'])
```

```
[11]: (array([ 35.,  1.,  2., 13., 107., 261., 243.,  87., 14.,  5.]),
      array([ 0., 12.2, 24.4, 36.6, 48.8, 61., 73.2, 85.4, 97.6,
            109.8, 122. ]),
      <BarContainer object of 10 artists>)
```




```
[12]: df[df['SkinThickness']==0].head()
```

```
[12]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
2           8      183             64              0         0  23.3
5           5      116             74              0         0  25.6
7          10      115              0              0         0  35.3
9           8      125             96              0         0   0.0
10          4      110             92              0         0  37.6
```

	DiabetesPedigreeFunction	Age	Outcome
2	0.672	32	1
5	0.201	30	0
7	0.134	29	0
9	0.232	54	1
10	0.191	30	0

```
[13]: plt.hist(df['SkinThickness'])
```

```
[13]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),
array([ 0. , 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),
<BarContainer object of 10 artists>)
```



```
[14]: df[df['Insulin']==0].head()
```

```
[14]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6     148           72           35         0  33.6
1           1      85           66           29         0  26.6
2           8     183           64            0         0  23.3
5           5     116           74            0         0  25.6
7          10     115            0            0         0  35.3
```

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1

```

5          0.201  30      0
7          0.134  29      0


```

```
[15]: plt.hist(df['Insulin'])
```

```

[15]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
      array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,
              761.4, 846. ]),
      <BarContainer object of 10 artists>)

```



```
[16]: df[df['BMI']==0].head()
```

```


[16]:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
9             8     125             96              0         0  0.0
49            7     105              0              0         0  0.0
60            2      84              0              0         0  0.0
81            2      74              0              0         0  0.0
145           0     102             75             23         0  0.0

   DiabetesPedigreeFunction  Age  Outcome
9                        0.232   54        1
49                        0.305   24        0
60                        0.304   21        0
81                        0.102   22        0
145                       0.572   21        0

```

```
[17]: plt.hist(df['BMI'])
```

```
[17]: (array([ 11.,  0.,  15., 156., 268., 224.,  78.,  12.,   3.,   1.]),  
      array([ 0. ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,  
            60.39, 67.1 ]),  
      <BarContainer object of 10 artists>)
```



inference: we clearly observe that column like BP, Glucose etc have 0 value. It isn't medically possible for some data record to have 0 value such as Blood Pressure or Glucose levels. Hence we replace them with the median value of that particular column.

```
[18]: #replacing 0 value with median
```

```
df['Glucose']=df['Glucose'].replace(0,df['Glucose'].median())  
df['BloodPressure']=df['BloodPressure'].replace(0,df['BloodPressure'].median())  
df['SkinThickness']=df['SkinThickness'].replace(0,df['SkinThickness'].median())  
df['Insulin']=df['Insulin'].replace(0,df['Insulin'].median())  
df['BMI']=df['BMI'].replace(0,df['BMI'].median())
```


(3) Create a count (frequency) plot describing the data types and the count of variables.

```
[19]: vc=df.dtypes.value_counts()
```



```
[20]: plt.figure(figsize=(8,6))
plt.bar(vc.index.astype(str),vc.values)
plt.xlabel('Data Type')
plt.ylabel('Frequency')
plt.title('Freq of data types')
```

```
[20]: Text(0.5, 1.0, 'Freq of data types')
```



2.2 Week 2:

2.2.1 Data Exploration:


(1) Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action:

```
[21]: target_count=df['Outcome'].value_counts()
target_count
```

```
[21]: 0    500  
      1    268  
      Name: Outcome, dtype: int64
```

```
[ ]:
```

```
[22]: # plot a counter plot to better understand our target feature  
plt.figure(figsize=(8,4))  
sns.countplot(x = 'Outcome',data = df)  
plt.xlabel('Target')  
plt.ylabel('Counts')  
plt.title('Target value of counts')  
plt.show()
```

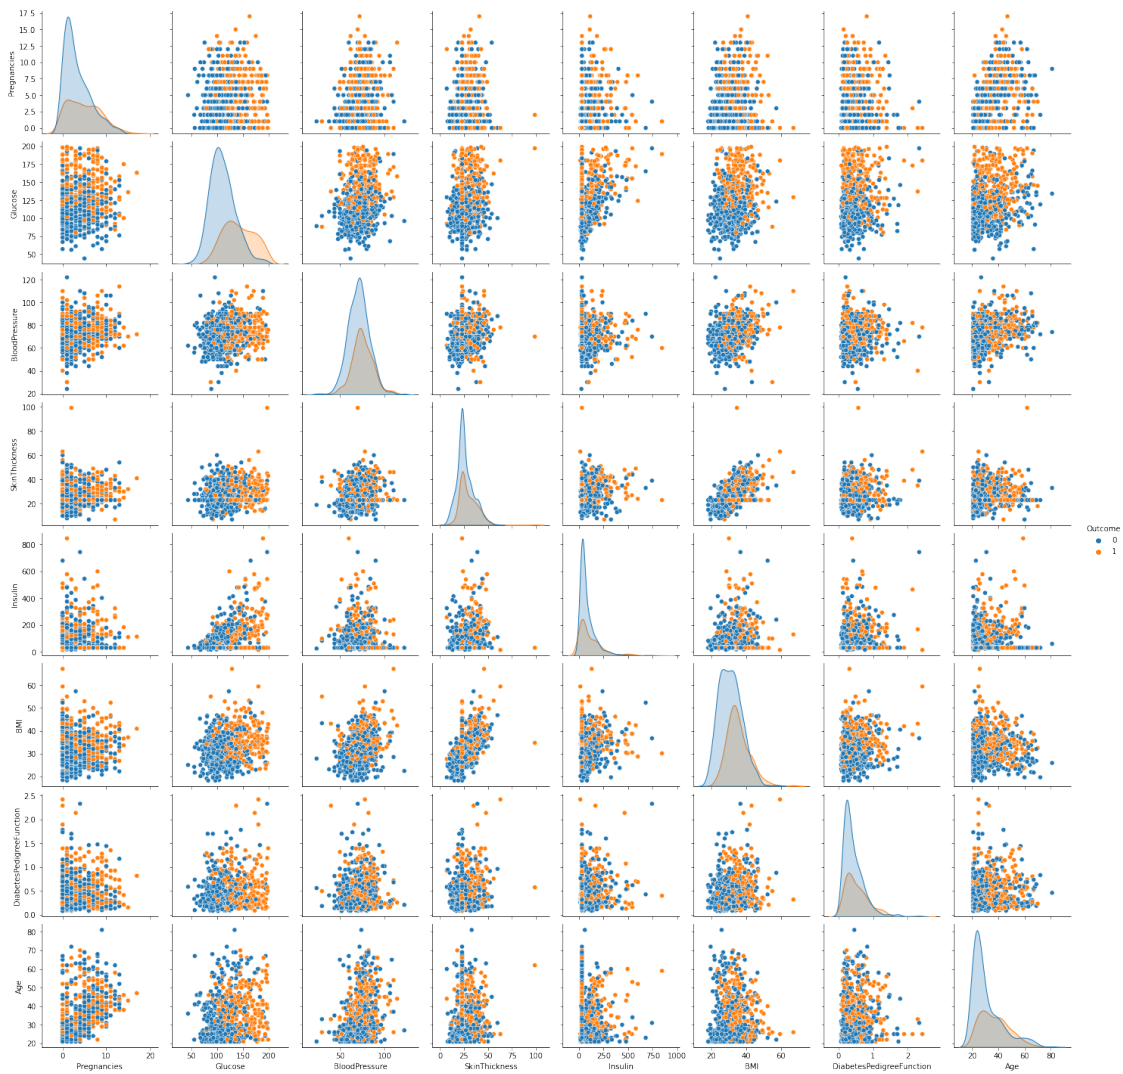


Inference: Our target variable has only two classes, 0 and 1. We can see that we have an imbalanced dataset. We have to make it balanced by using sampling techniques to avoid overfitting.

(2) Create scatter charts between the pair of variables to understand the relationships. Describe your findings:


```
[23]: #creating pair plot  
import seaborn as sns  
sns.pairplot(df,hue='Outcome')
```

```
[23]: <seaborn.axisgrid.PairGrid at 0x7fe5b3de1e10>
```




```
[24]: #let's find the relationship between independent and dependent numerical
      ↪ features
      sns.scatterplot(x='Glucose',y='BloodPressure',data=df,hue='Outcome')
```

```
[24]: <AxesSubplot:xlabel='Glucose', ylabel='BloodPressure'>
```




```
[25]: sns.scatterplot(y='Insulin',x='Glucose',data=df,hue='Outcome')
```

```
[25]: <AxesSubplot:xlabel='Glucose', ylabel='Insulin'>
```



```
[26]: sns.scatterplot(y='BMI',x='Age',data=df,hue='Outcome')
```

```
[26]: <AxesSubplot:xlabel='Age', ylabel='BMI'>
```



We have some observations from above scatter plot of pairs of features: * **Glucose** alone is impressively good to distinguish between the **Outcome** classes. * **Age** is also able to distinguish between classes to some extent

(3) Perform correlation analysis. Visually explore it using a heat map:

```
[27]: # Finding the correlation between each features
df.corr()
```

```
[27]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.128213	0.208615	0.032568	
Glucose	0.128213	1.000000	0.218937	0.172143	
BloodPressure	0.208615	0.218937	1.000000	0.147809	
SkinThickness	0.032568	0.172143	0.147809	1.000000	
Insulin	-0.055697	0.357573	-0.028721	0.238188	
BMI	0.021546	0.231400	0.281132	0.546951	
DiabetesPedigreeFunction	-0.033523	0.137327	-0.002378	0.142977	
Age	0.544341	0.266909	0.324915	0.054514	
Outcome	0.221898	0.492782	0.165723	0.189065	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.055697	0.021546	-0.033523	
Glucose	0.357573	0.231400	0.137327	
BloodPressure	-0.028721	0.281132	-0.002378	
SkinThickness	0.238188	0.546951	0.142977	
Insulin	1.000000	0.189022	0.178029	
BMI	0.189022	1.000000	0.153506	
DiabetesPedigreeFunction	0.178029	0.153506	1.000000	
Age	-0.015413	0.025744	0.033561	
Outcome	0.148457	0.312249	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.266909	0.492782
BloodPressure	0.324915	0.165723
SkinThickness	0.054514	0.189065
Insulin	-0.015413	0.148457
BMI	0.025744	0.312249
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[28]: plt.figure(figsize=(12,8))
      sns.heatmap(df.corr(),annot=True,cmap='YlGnBu')
```

```
[28]: <AxesSubplot:>
```



It appears from correlation matrix and heatmap that there exists significant correlation between some pairs such as - * Age-Pregnancies * BMI-SkinThickness

2.3 Week 3:

2.3.1 Data Modeling:

(1) Devise strategies for model building. It is important to decide the right validation framework. Express your thought process:

Answer: Since this is a classification problem, we will be building popular classification models for our training data and then compare performance of each model on test data to accurately predict target variable (Outcome):

- 1) Logistic Regression
- 2) Decision Tree
- 3) RandomForest Classifier
- 4) K-Nearest Neighbour (KNN)

```
[29]: #splitting ind and dep var  
X=df.drop(columns=['Outcome'])  
y=df.Outcome
```

```
[30]: ! pip install imblearn
```

```
Defaulting to user installation because normal site-packages is not writeable  
Requirement already satisfied: imblearn in ./local/lib/python3.7/site-packages  
(0.0)  
Requirement already satisfied: imbalanced-learn in  
/usr/local/lib/python3.7/site-packages (from imblearn) (0.8.0)  
Requirement already satisfied: scikit-learn>=0.24 in  
/usr/local/lib/python3.7/site-packages (from imbalanced-learn->imblearn) (1.0.2)  
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/site-  
packages (from imbalanced-learn->imblearn) (0.14.1)  
Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.7/site-  
packages (from imbalanced-learn->imblearn) (1.4.1)  
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.7/site-  
packages (from imbalanced-learn->imblearn) (1.21.5)  
Requirement already satisfied: threadpoolctl>=2.0.0 in  
/usr/local/lib/python3.7/site-packages (from scikit-learn>=0.24->imbalanced-  
learn->imblearn) (2.2.0)  
WARNING: You are using pip version 22.0.3; however, version 23.1.2 is  
available.  
  
You should consider upgrading via the '/usr/local/bin/python3 -m pip install  
--upgrade pip' command.
```

```
[31]: from imblearn.over_sampling import SMOTE  
smk = SMOTE()  
x_train_smote,y_train_smote=smk.fit_resample(X,y)
```

```
[32]: from collections import Counter  
print('Original dataset shape {}'.format(Counter(y)))  
print('Resampled dataset shape {}'.format(Counter(y_train_smote)))
```

```
Original dataset shape Counter({0: 500, 1: 268})  
Resampled dataset shape Counter({1: 500, 0: 500})
```

```
[33]: print(x_train_smote.shape)  
print(y_train_smote.shape)
```

```
(1000, 8)  
(1000,)
```



```
[34]: # Split the data set into training and testing
      from sklearn.model_selection import train_test_split

      x_train, x_test, y_train, y_test = \
      ↪train_test_split(x_train_smote, y_train_smote, test_size=0.2, random_state=3)
```

```
[35]: from sklearn.preprocessing import StandardScaler

      scaler = StandardScaler()
      x_train_std = scaler.fit_transform(x_train)
      x_test_std = scaler.transform(x_test)
```

```
[36]: x_train_std.shape
```

```
[36]: (800, 8)
```

```
[37]: x_test_std.shape
```

```
[37]: (200, 8)
```

2.3.2 2. Apply an appropriate classification algorithm to build a model. Compare various models with the results from KNN algorithm.

```
[38]: # Building Model For Logistic Regression
      from sklearn.linear_model import LogisticRegression
      lr_model = LogisticRegression()
```

```
[39]: # Training Logistic Regression Model
      # Logistic regression works well on scaled
      lr_model.fit(x_train_std, y_train)
      y_predict = lr_model.predict(x_test_std)
```

```
[40]: from sklearn.metrics import confusion_matrix
      from sklearn.metrics import accuracy_score, recall_score, f1_score
      from sklearn.metrics import roc_auc_score, classification_report, roc_curve
```

```
[41]: # Getting all accuracy scores for Logistic Regression
      accuracy = accuracy_score(y_test, y_predict)
      recall = recall_score(y_test, y_predict)
      f1 = f1_score(y_test, y_predict)

      print("Accuracy: ", accuracy)
      print("Recall: ", recall)
      print("F1: ", f1)
```

Accuracy: 0.755
Recall: 0.7289719626168224
F1: 0.7609756097560975

```
[42]: # Building Model For Decision Tree classifier
from sklearn.tree import DecisionTreeClassifier
dt_model = DecisionTreeClassifier(max_depth=5, random_state=42)
# scaling does affect decision tree
dt_model.fit(x_train_std, y_train)
y_pred = dt_model.predict(x_test_std)
```

```
[43]: # getting all types of accuracy for decision tree
dt_accuracy = accuracy_score(y_test,y_pred)
dt_recall = recall_score(y_test,y_pred)
dt_f1 = f1_score(y_test,y_pred)

print("Accuracy: ", dt_accuracy)
print("Recall: ", dt_recall)
print("F1: ", dt_f1)
```

Accuracy: 0.755
Recall: 0.8130841121495327
F1: 0.7802690582959642


```
[44]: # Building model for Random Forest Classifier
from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier(n_estimators=100,max_depth=5,random_state=42)
rf_model.fit(x_train_std, y_train)
y_pred1 = rf_model.predict(x_test_std)
```

```
[45]: # getting all types of evaluation scores
rf_accuracy = accuracy_score(y_test,y_pred1)
rf_recall = recall_score(y_test,y_pred1)
rf_f1 = f1_score(y_test,y_pred1)

print("Accuracy: ", rf_accuracy)
print("Recall: ", rf_recall)
print("F1: ", rf_f1)
```

Accuracy: 0.835
Recall: 0.8598130841121495
F1: 0.847926267281106

```
[46]: # plotting confusion matrix for random forest
rf_cm = confusion_matrix(y_test, y_pred1)
sns.heatmap(rf_cm, annot=True)
plt.show()
```



```
[47]: # Printing overall report for random forest
print(classification_report(y_test, y_pred1))
```

	precision	recall	f1-score	support
0	0.83	0.81	0.82	93
1	0.84	0.86	0.85	107
accuracy			0.83	200
macro avg	0.83	0.83	0.83	200
weighted avg	0.83	0.83	0.83	200

```
[48]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=20).fit(x_train_std,y_train)
y_pred2 = knn.predict(x_test_std)
```

```
[49]: # getting all types of evaluation scores
knn_accuracy = accuracy_score(y_test,y_pred2)
knn_recall = recall_score(y_test,y_pred2)
knn_f1 = f1_score(y_test,y_pred2)

print("Accuracy: ", knn_accuracy)
print("Recall: ", knn_recall)
print("F1: ", knn_f1)
```

Accuracy: 0.775
Recall: 0.8130841121495327
F1: 0.7945205479452055


Among all models, RandomForest has given best accuracy, recall and f1_score. Therefore we will build final model using RandomForest.

2.4 Week 4:

2.4.1 Data Modeling:

(1) Create a classification report by analyzing sensitivity, specificity, AUC (ROC curve), etc. Please be descriptive to explain what values of these parameter you have used:

```
[50]: # plotting confusion metrix for random forest
rf_cm = confusion_matrix(y_test, y_pred1)
sns.heatmap(rf_cm, annot=True)
plt.show()
```



```
[51]: # Printing overall report for random forest
print(classification_report(y_test, y_pred1))
```


	precision	recall	f1-score	support
0	0.83	0.81	0.82	93

	1	0.84	0.86	0.85	107
accuracy				0.83	200
macro avg		0.83	0.83	0.83	200
weighted avg		0.83	0.83	0.83	200

```
[52]: fpr,tpr,thresholds=roc_curve(y_test,y_pred1)
```

```
[53]: auc=roc_auc_score(y_test,y_pred1)
```

```
[54]: plt.figure(figsize=(8,6))
plt.plot(fpr,tpr,label='ROC Curve (AUC={:.2f})'.format(auc))
plt.plot([0,1],[0,1],linestyle='--',color='r',label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



In this scenario, we encountered a data imbalance issue, where the distribution of classes in the dataset was not uniform. To address this problem, we applied the Synthetic Minority Over-sampling

Technique (SMOTE) to balance the data. The objective was to improve the performance of a classification model by ensuring equal representation of the minority class.

After applying SMOTE, we evaluated the model's performance using several metrics. The accuracy of the model was measured to be 83%, indicating that 83% of the predictions made by the model were correct. This metric provides an overall assessment of the model's performance.

Additionally, we assessed the model's ability to correctly identify positive instances of the minority class. The recall, also known as sensitivity or true positive rate, was measured to be 85%. This means that out of all the actual positive instances, the model successfully identified 85% of them. A higher recall score indicates a better ability to detect positive cases.

Furthermore, we examined the model's balance between precision and recall using the F1 score. The F1 score, which is the harmonic mean of precision and recall, was calculated to be 84%. This score provides a balanced assessment of the model's performance in terms of both false positives and false negatives. A higher F1 score indicates a better trade-off between precision and recall.

Finally, we assessed the model's ability to rank instances and discriminate between positive and negative classes using the Receiver Operating Characteristic (ROC) curve. The area under the ROC curve (ROC AUC score) was calculated to be 0.83. This metric measures the model's ability to distinguish between positive and negative instances. A higher ROC AUC score indicates better discrimination ability.

In conclusion, the application of SMOTE to address the data imbalance issue improved the performance of the classification model. The model achieved an accuracy of 83%, recall of 85%, F1 score of 84%, and an ROC AUC score of 0.83. These results indicate that the model has demonstrated promising performance in handling the imbalanced data and is effective in identifying positive instances of the minority class. However, it is important to note that further evaluation and validation on additional datasets may be necessary to ascertain the generalizability of these results.

2.4.2 Data Reporting:

2. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:

- a. Pie chart to describe the diabetic or non-diabetic population
- b. Scatter charts between relevant variables to analyze the relationships
- c. Histogram or frequency charts to analyze the distribution of the data
- d. Heatmap of correlation analysis among the relevant variables
- e. Create bins of these age values: 20-25, 25-30, 30-35, etc. Analyze different variables :

2.5 PLEASE REFER TABLEAU PUBLIC LINK FOR DASHBOARD AND VISUALIZATION.

[]: